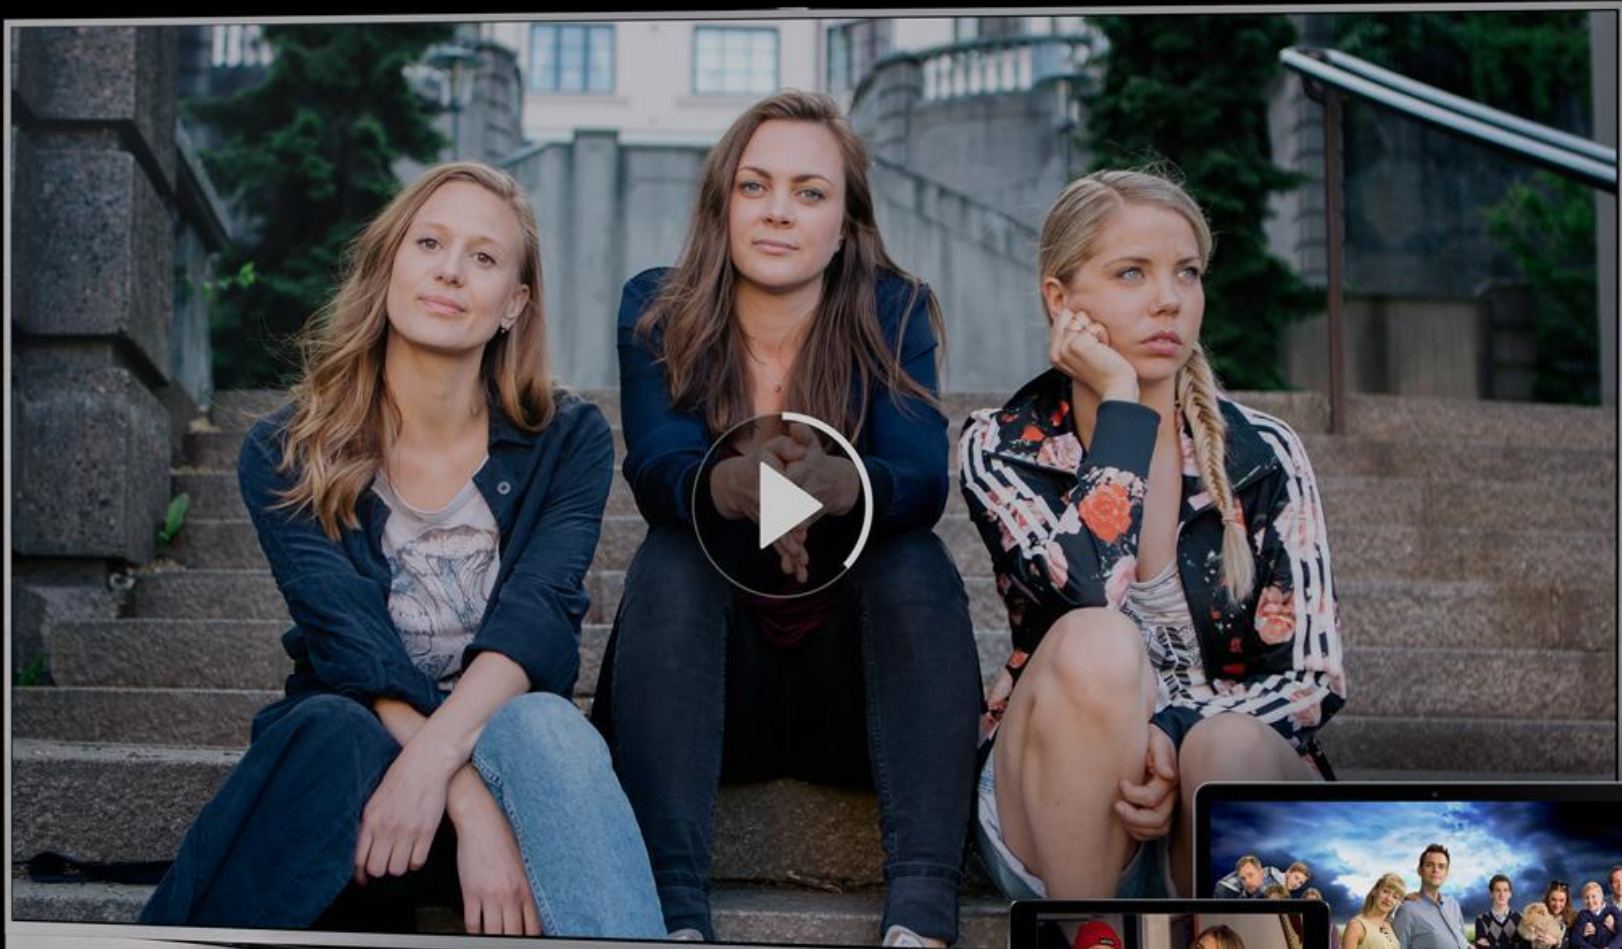


**BUILDING
RELIABLE SERVICES
AT NRK TV**

@EINARWH

NIK TV



AGENDA

WHAT IS THE SERVICE?
WHAT IS RELIABILITY?
RELIABILITY ENGINEERING
STRESS TESTS
TOOLS & TECHNIQUES

WHAT IS THE SERVICE?

WHAT IS RELIABILITY?

RELIABILITY ENGINEERING

STRESS TESTS

TOOLS & TECHNIQUES

WHAT IS THE SERVICE?

WHAT IS RELIABILITY?

RELIABILITY ENGINEERING

STRESS TESTS

TOOLS & TECHNIQUES

WHAT IS THE SERVICE?

WHAT IS RELIABILITY?

RELIABILITY ENGINEERING

STRESS TESTS

TOOLS & TECHNIQUES

WHAT IS THE SERVICE?
WHAT IS RELIABILITY?
RELIABILITY ENGINEERING
STRESS TESTS
TOOLS & TECHNIQUES

WHAT IS THE SERVICE?
WHAT IS RELIABILITY?
RELIABILITY ENGINEERING
STRESS TESTS
TOOLS & TECHNIQUES

WHAT IS THE SERVICE?

TV STREAMING SERVICE

TV STREAMING APP

CLIENT - SERVER

PLAY THE GAME OF HUNT THE MEDIA MANIFEST

CONTENT DELIVERY NETWORK **FOR THE PETABYTES**

OUR API IS A METADATA API

ALMOST EXCLUSIVELY READS

WRITES FOR PROGRESS AND ANALYTICS

TRAFFIC BURSTS

BIG LAUNCH = BLACK FRIDAY

WHAT IS A RELIABLE SERVICE?

**WHAT DOES IT MEAN
TO BE RELIABLE?**

ROBUST AND RESILIENT

**WHAT DOES IT MEAN
TO BE ROBUST?**

WHAT DOES IT MEAN TO BE **RESILIENT**?

ROBUST



* CROWD CHEERING *

RESILIENT

Tubthumping

Chumbawamba

I get knocked down, but I get up again

You are never gonna keep me down

I get knocked down, but I get up again

You are never gonna keep me down

I get knocked down, but I get up again

You are never gonna keep me down

I get knocked down, but I get up again

You are never gonna keep me down



**ROBUST AND RESILIENT
ARE NOT BOOLEANS**

ROBUST AND RESILIENT
ARE NOT EVEN PERCENTAGES

**ROBUST AGAINST
CERTAIN KINDS AND
AMOUNTS OF STRESS**

RESILIENT UNDER CERTAIN KINDS OF BAD CONDITIONS

NON-FUNCTIONAL REQUIREMENTS

SPECIFIC AND CONTEXTUAL

DEPENDS ON THE SERVICE

BUSINESS REQUIREMENTS

«GOOD ENOUGH»

TYPICALLY NOT SPECIFIED

«BEST EFFORT»

BEST EFFORT IS A GIVEN

BEST EFFORT
ISN'T AUTOMATICALLY
GOOD ENOUGH

SOMETIMES
GOOD ENOUGH
MUST BE
REALLY QUITE GOOD

**WE MIGHT NOT
REACH IT THROUGH
BEST EFFORT ALONE**

MINDSET

**MOST DEVELOPERS ARE
CONSCIENTIOUS**

COMPILE TIME MINDSET

«THE CODE IS THE TRUTH»

WHAT TRUTH?

THE TRUTH ABOUT WHAT?

WHAT HAPPENS IN PRODUCTION?

RUNTIME ENVIRONMENT

AVAILABLE RESOURCES

EXTERNAL DEPENDENCIES

**THE RUNTIME IS
DYNAMIC**

**THE RUNTIME IS PARTIALLY
OUT OF YOUR CONTROL**

**THE RUNTIME IS
UNKNOWNABLE**

**RUNTIME BEHAVIOR IS
NOT PURELY ANALYTICAL**

CODE + ENVIRONMENT

RUNTIME MINDSET

**WE CAN'T FIGURE IT
ALL OUT IN OUR HEADS**

SYSTEMS THINKING

HYPOTHESES AND VALIDATION

BEING ON-CALL **DOES WONDERS FOR** **THE RUNTIME MINDSET**

HOW COMPLEX SYSTEMS FAIL

RICHARD COOK

RELIABILITY AT NRK TV

IMPROVEMENT START WITH A CRISIS

**OUR BEST EFFORT
WASN'T GOOD ENOUGH**

HOW TO IMPROVE BEYOND BEST EFFORT?

**HOW CAN WE TELL IF IT IS
GETTING BETTER OR WORSE?**

WE NEED FEEDBACK

ENGINEERING

=

BEST EFFORT + FEEDBACK

FROM **AD HOC** TO **STRATEGIC** RELIABILITY WORK

SERVICE LEVEL OBJECTIVE

[SLO]

SERVICE LEVEL INDICATOR

[SLI]

ERROR RATE

LATENCY

THROUGHPUT

CAVEATS

METRICS CAN BE HARMFUL

METRICS CAN BE ABUSED

METRICS CAN BE GAMED

UNDERMINES THE PURPOSE

REQUIRES TRUST

BOTTOM-UP

SELF-IMPOSED BY TEAM

DEFINED BY TEAM

OWNED BY TEAM

NO REWARDS

NO COMPETITION

**AN SLO SHOULD BE
APPROPRIATE**

**INVESTING IN MEETING A
TOO-STRICT SLO IS WASTE**

NO VANITY METRICS

WHAT ARE SLOS FOR?

GUIDE DECISIONS

TO HELP COMMUNICATE

TO VALIDATE IMPROVEMENTS

TO AVOID REGRESSIONS

HOW TO DEFINE AN SLO

UNCOMFORTABLE!

UNCOMFORTABLE!
(IT MUST BE GOOD!)

FORCE DECISIONS

MAKE IT CONCRETE

SLO Channel Playback endpoints

Uptime: 99.95%

Measured behind gateway/cache.

Measuring interval: 1 minute.

SLO applies to loads of 10 to 50.000 req/min.

SLO applies in time interval 06:00 - 01:00.

	5xx errors	Latency (99 percentile)
PS Playback API Channel - tv/live	< 0.01 %	< 200 ms
PS Playback API Channel - playback/channel/metadata	< 0.01 %	< 100 ms
PS Playback API Channel - playback/channel/manifest	< 0.01 %	< 200 ms

START MEASURING
CHOOSE OBJECTIVES
OBSERVE
REVISE/INVEST

=> START MEASURING <=

CHOOSE OBJECTIVES

OBSERVE

REVISE/INVEST

START MEASURING

=> CHOOSE OBJECTIVES <=

OBSERVE

REVISE/INVEST

START MEASURING
CHOOSE OBJECTIVES
=> OBSERVE <=
REVISE/INVEST

START MEASURING
CHOOSE OBJECTIVES
OBSERVE
=> REVISE/INVEST <=

START MEASURING
CHOOSE OBJECTIVES
=> OBSERVE <=
REVISE/INVEST

START MEASURING
CHOOSE OBJECTIVES
OBSERVE
=> REVISE/INVEST <=

OBSERVATION

DASHBOARDS

```

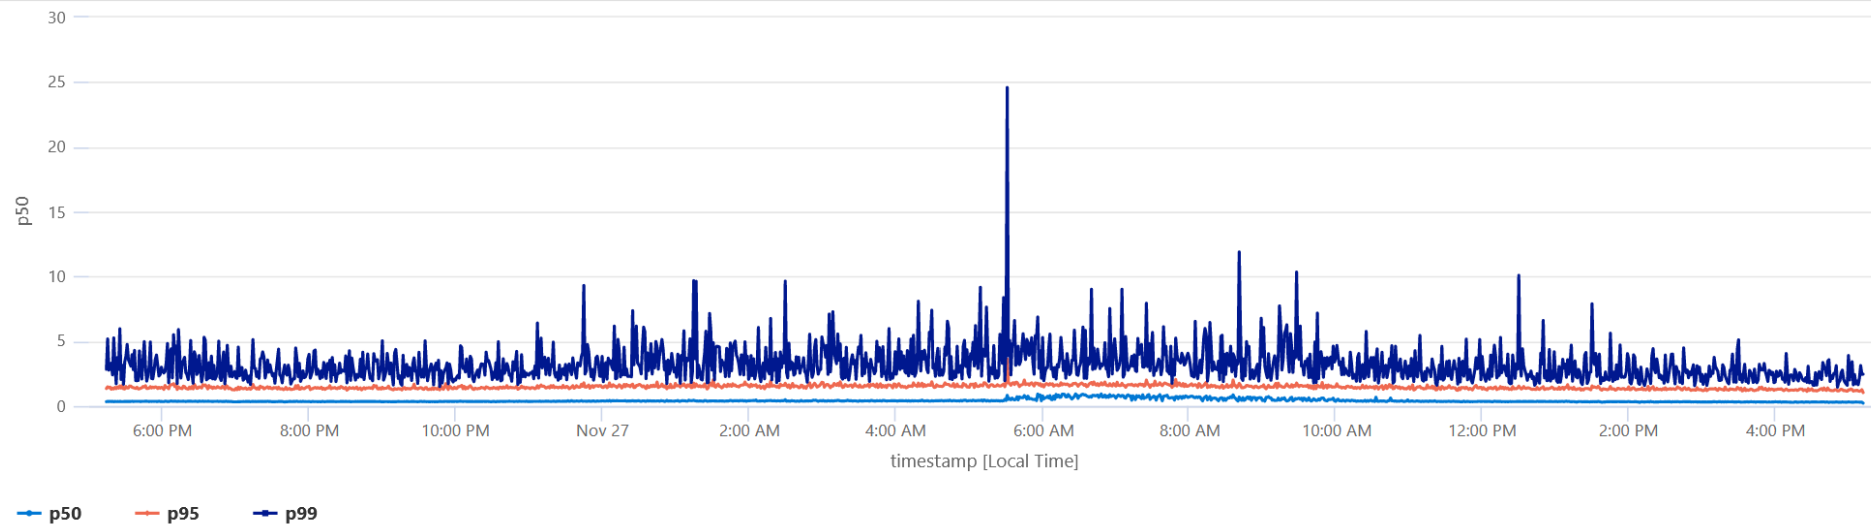
1 requests
2 | where customDimensions.applicationName == "ps-playback-api-channel"
3 | where url contains "metadata/channel"
4 | summarize percentiles(duration, 50, 95, 99) by bin(timestamp, 1m)
5 | project
6 |     timestamp,
7 |     p50=percentile_duration_50,
8 |     p95=percentile_duration_95,
9 |     p99=percentile_duration_99
10 | render timechart
11

```

Results Chart | ⌚ Display time (UTC+01:00) ▾

Completed. Showing results from the last 24 hours.

⌚ 00:01.2 📄 1,439 records



REPORTS

Weekly SLO Uptime report for playback/manifest/channel

SLO report for the week 11/15/2021 12:00:08 AM - 11/22/2021 12:00:08 AM

SLO latency 99% pr min: < 200 ms

SLO error rate pr min: < 0.01%

SLO applies to the time interval 06:00 - 01:00

SLO uptime: 99.95%

SLO downtime budget: 4 minutes.

SLO Status: ✓

Uptime: 99.99%

Downtime: 1 minutes.

Downtime (latency): 1 minutes.

Downtime (error rate): 0 minutes.

Additional downtime outside SLO hours: 1 minutes.

NOTIFICATIONS?

**SHOULD WE SET UP ALERTS
FOR SLO VIOLATIONS?**

**ALERTS SHOULD MEAN
ACTION NECESSARY**

IT'S NOT A HEALTH CHECK

ARE WE RELIABLE NOW?

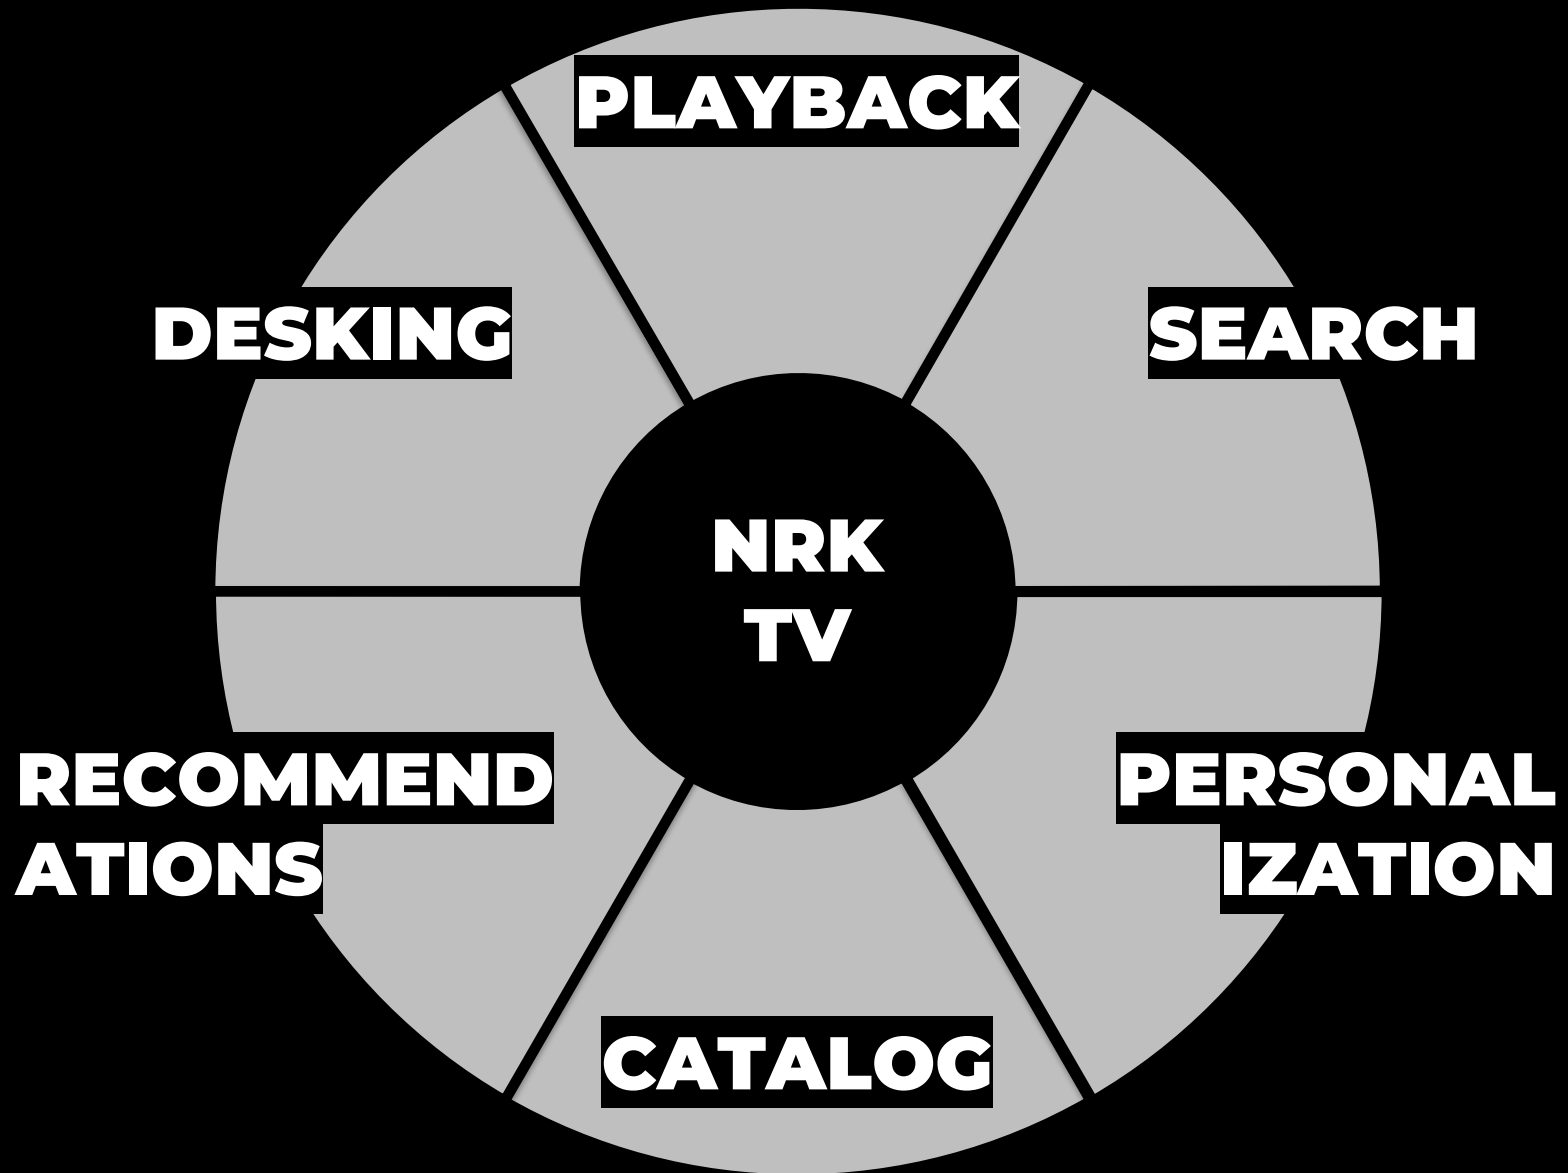
OF COURSE NOT!

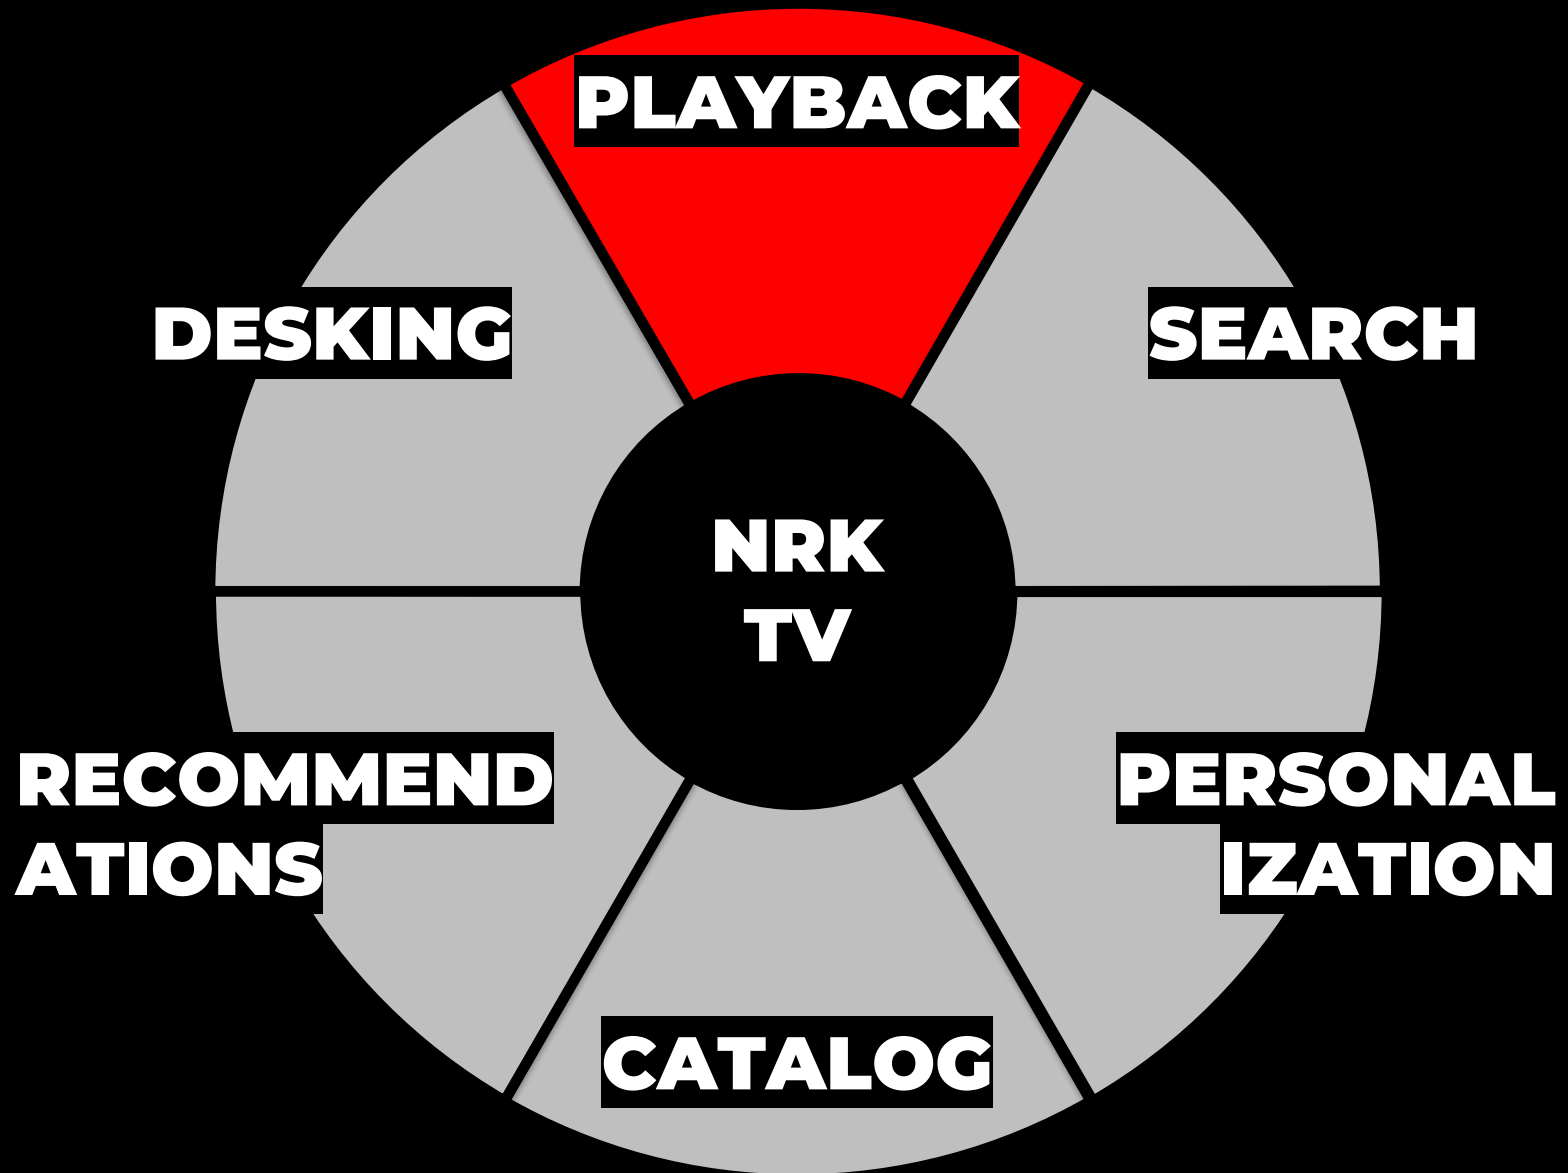
WE HAVE FEEDBACK

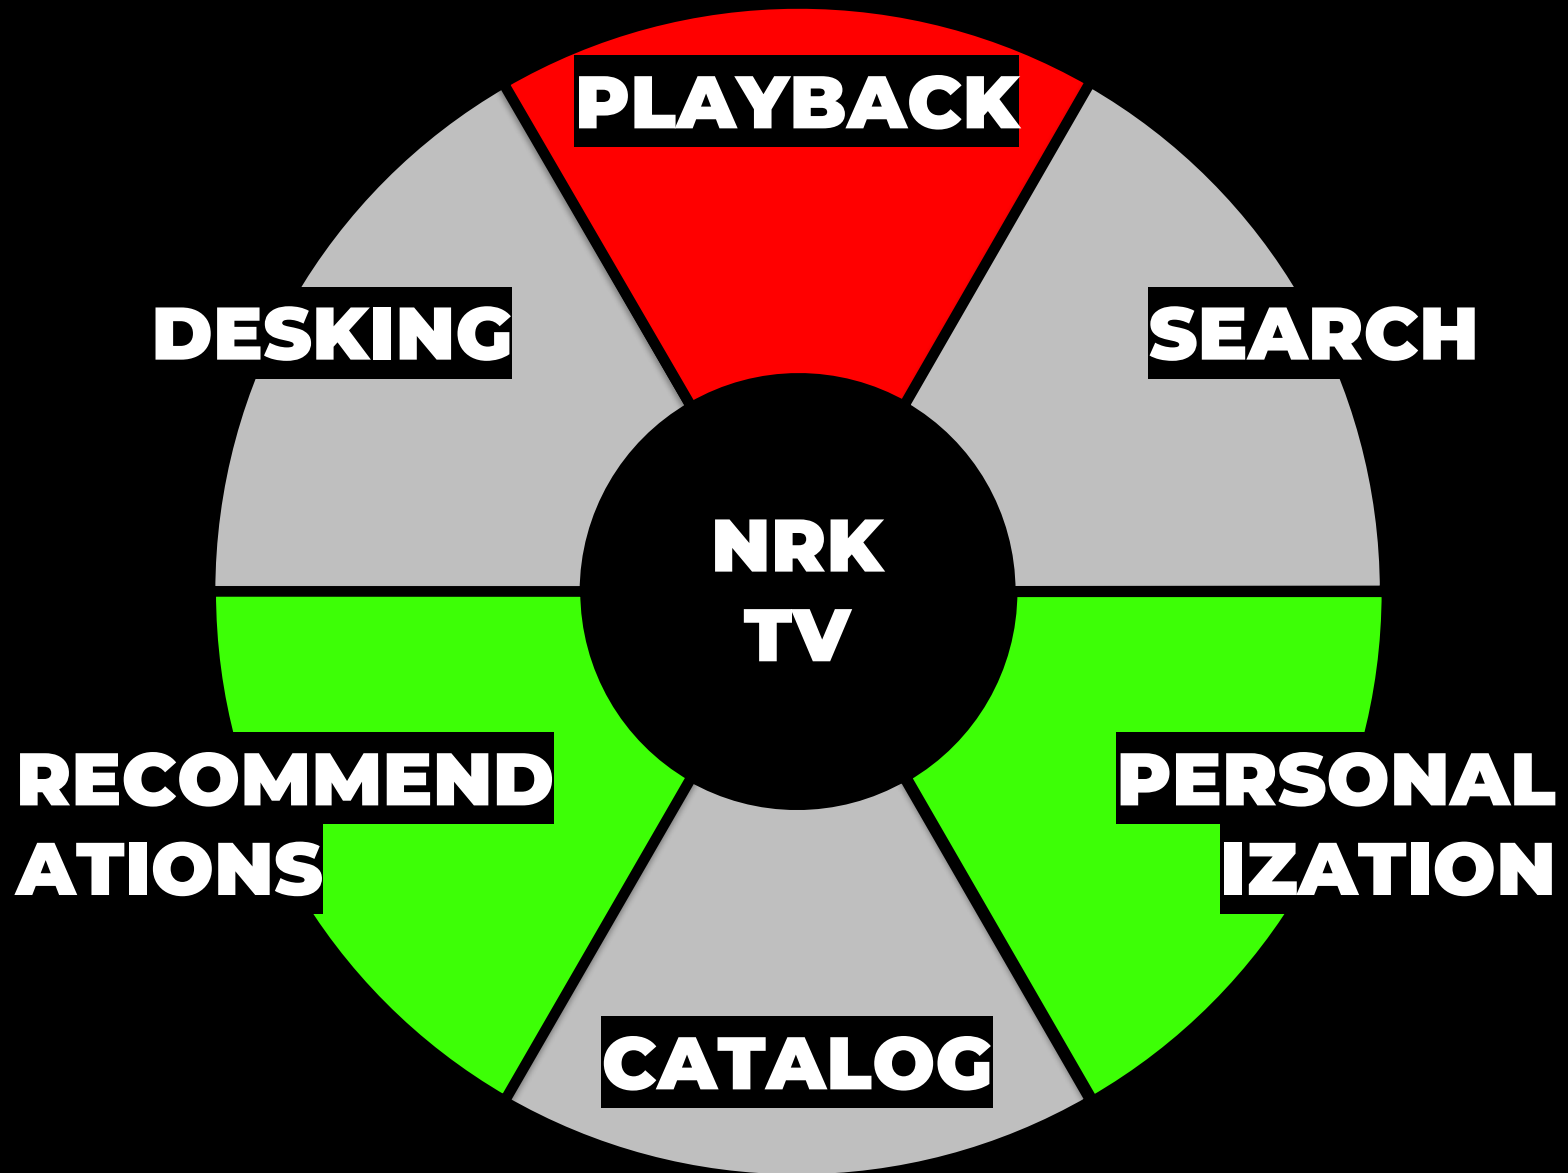
WE CAN BEGIN WORKING

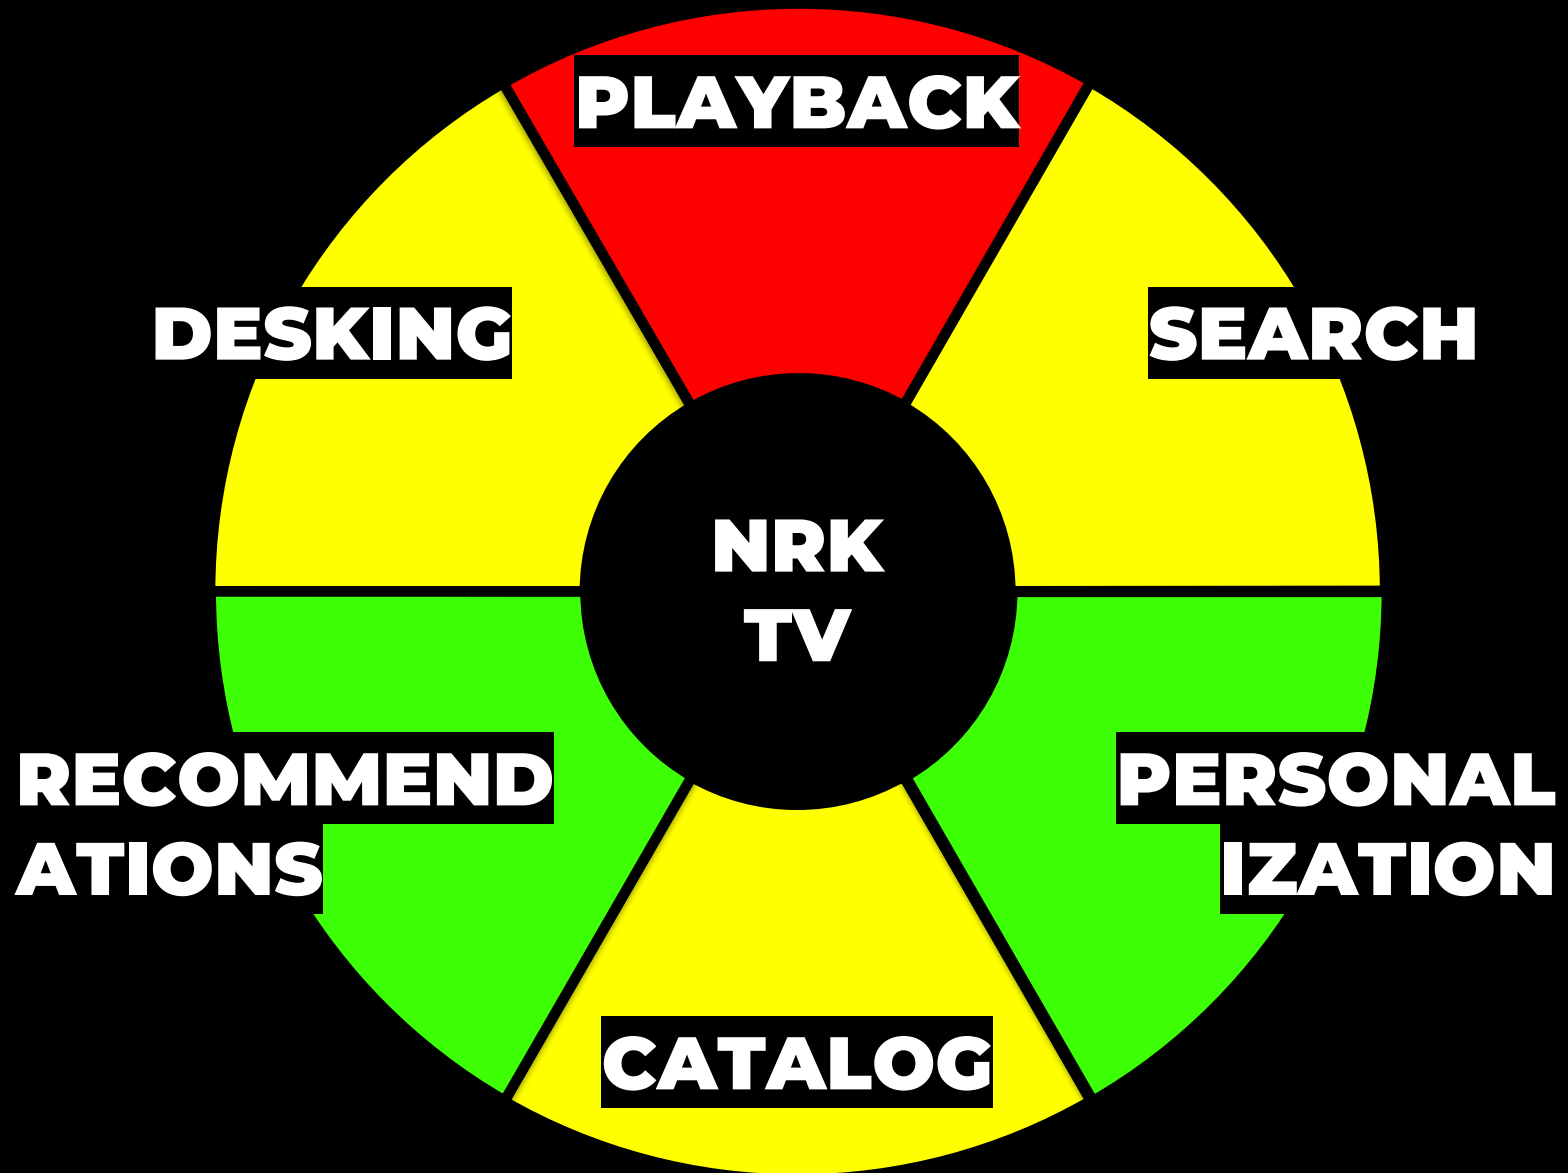
ARCHITECTING FOR RELIABILITY

RED AND GREEN SERVICES









SEPARATE SLOS

SEPARATE RELIABILITY MEASURES

INDEPENDENT SERVICES

FAILURE ISOLATION

PARTIAL FAILURE

PARTIAL SUCCESS

GRACEFUL DEGRADATION

VALIDATION

BROKEN PROXY

MOBILE APP



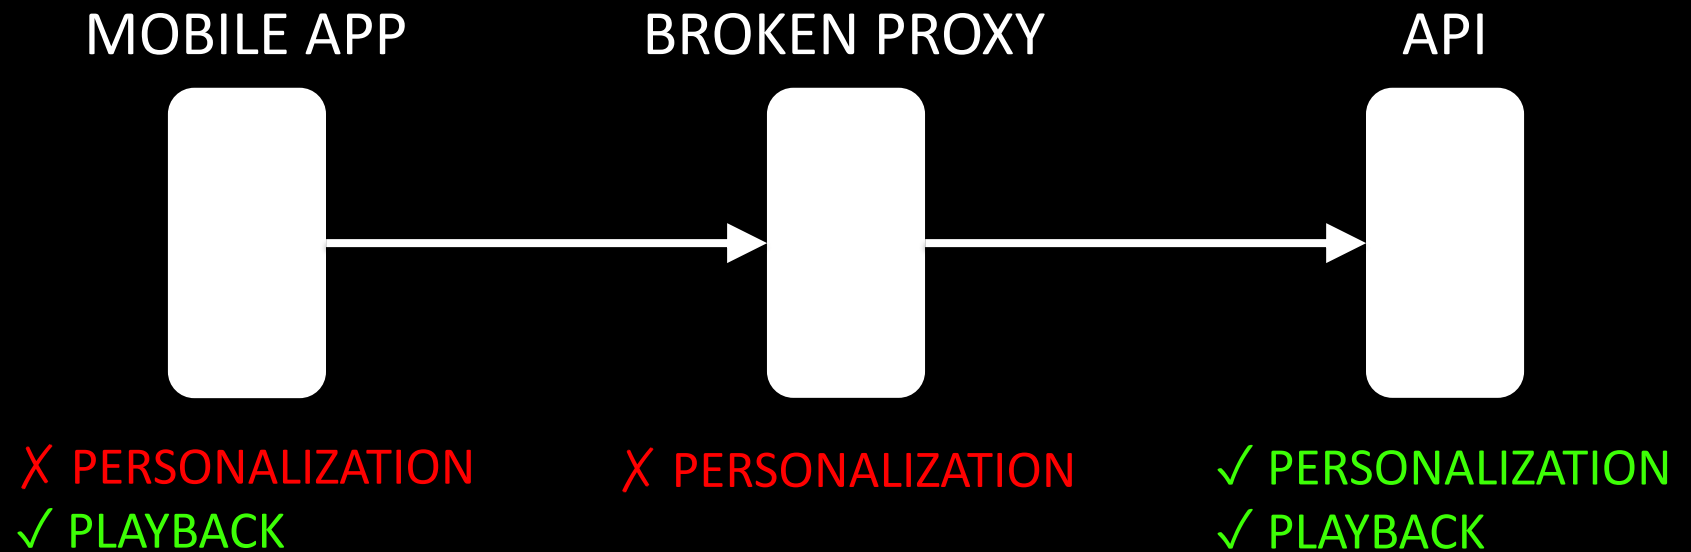
✓ PERSONALIZATION
✓ PLAYBACK

API



✓ PERSONALIZATION
✓ PLAYBACK

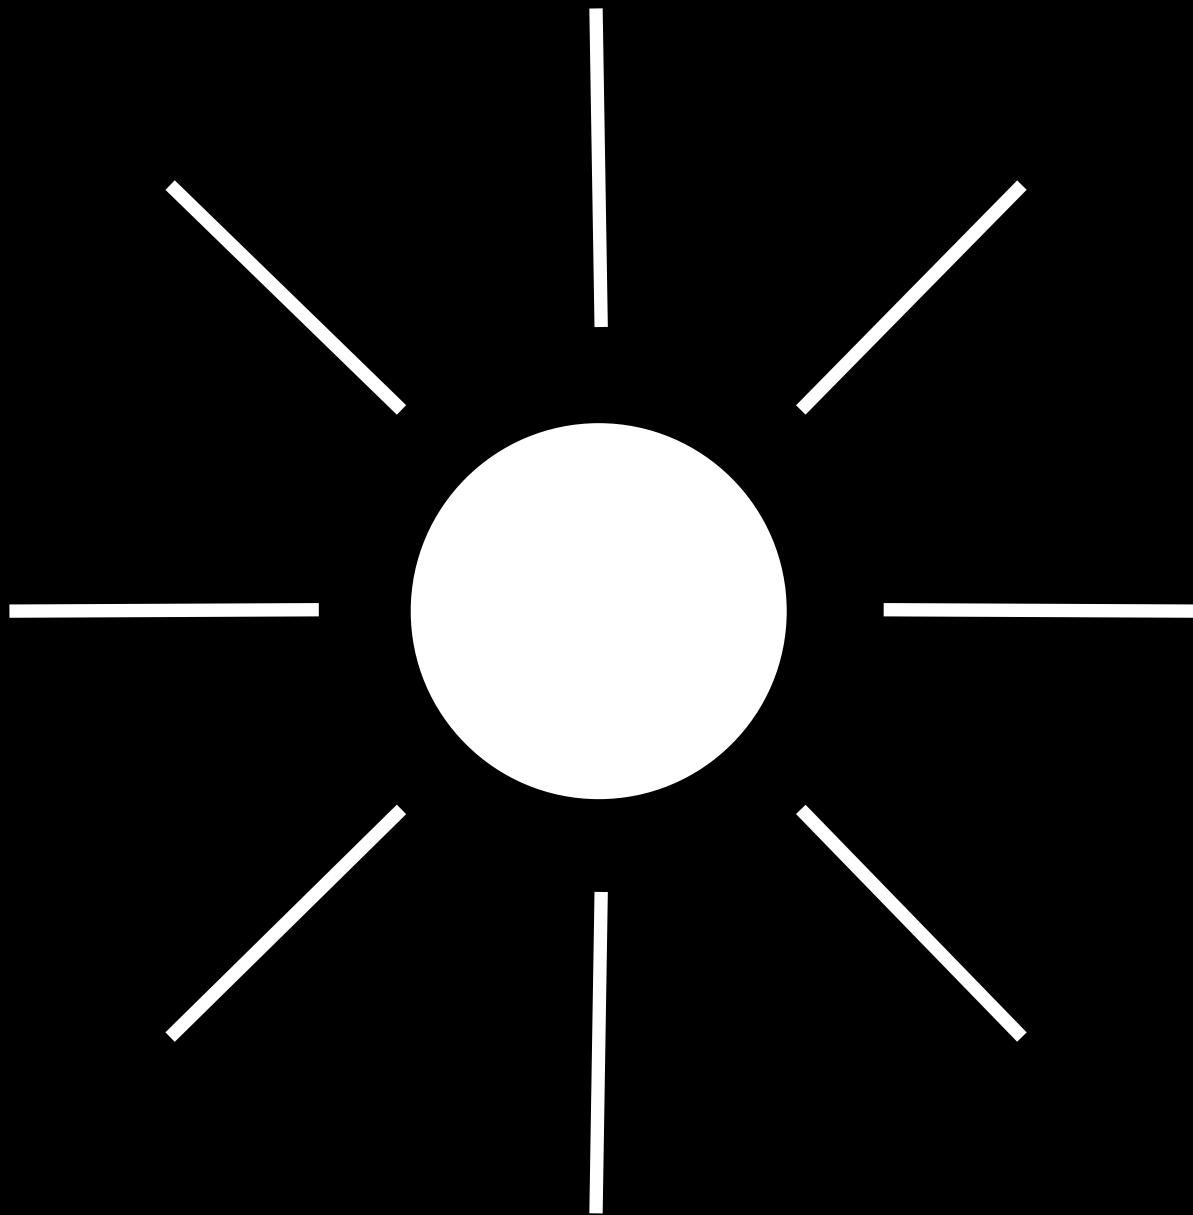




THIS IS RESILIENCE

WHAT ABOUT ROBUSTNESS?

MEETING SLO



MEETING SLO ON A BAD DAY

MEETING SLO UNDER STRESS

EXPERIMENTS

STRESS TESTS

TOOLS

LOAD TESTING



RESILIENCE LIBRARY



FAULT INJECTION LIBRARY



FEATURE TOGGING



PURE LOAD TESTS

VERIFY SCALABILITY

AUTOSCALE?

LOAD TESTS WITH INTRODUCED FAULTS

SIMULATE

FAILURE

WHAT CAN FAIL?

FAILING DEPENDENCY

HTTP CALL

DATABASE

SECONDARY FAILURES

EXHAUST LIMITED RESOURCES

HOW LONG DOES IT TAKE?

ASSUMPTION

FAST FAILURE

REALITY

SLOW FAILURE

STRATEGIES & REMEDIES

WHAT TO DO ABOUT FAILURE?

RELIABILITY PATTERNS

TIMEOUT

RETRY

CIRCUIT BREAKER



SOME QUESTIONS

DO WE NEED THE RESULT?

HOW BADLY DO WE NEED IT?

DEFAULT VALUE?

CACHED VALUE?

CACHING

HTTP CACHE

DESIGNING FOR CACHEABILITY

BÅRD'S TIME EXPANSION ALGORITHM

BÅRD'S

TIME INDEPENDENCE

PRINCIPLE

TIME-DEPENDENT TITLES

**REFRESH TO SEE
IF THE TITLE CHANGED?**

THIS IS STUPID

**TITLES CHANGE
VERY RARELY**

**THIS IS COMPLETELY
DETERMINISTIC**

GENERATE A TABLE

FROM	TO	TITLE
03.12.2021 00:00:00	04.12.2021 00:00:00	I dag
04.12.2021 00:00:00	05.12.2021 00:00:00	I går
05.12.2021 00:00:00	08.12.2021 00:00:00	Fredag
		3. desember

TIME-DEPENDENT TITLES
NO LONGER INFLUENCE
CACHEABILITY

WHAT TO DO ABOUT FAILURE?

HYPOTHESIS

VALIDATION

CASE

LINEAR CHANNEL

PLAYBACK API

WHY IS THE LINEAR CHANNEL PLAYBACK API IMPORTANT?

LOAD TESTS WITH INJECTED FAILURES

INJECTING FAILURE WITH SIMMY

CONTROLLING FAILURE CONFIG WITH UNLEASH

Search

channel-blob-storage-inject-fault



Last minute ▼

By name ▼



ps-playback-api-channel-blob-storage-inject-fault-always

Inject fault in all calls to blob storage for the channels.json blob using Simmy.



ps-playback-api-channel-blob-storage-inject-fault-low-rate

Inject fault in calls to blob storage for the channels.json blob using Simmy (low injection rate).



ps-playback-api-channel-blob-storage-inject-fault-moderate-rate

Inject fault in calls to blob storage for the channels.json blob using Simmy (moderate injection rate).



ps-playback-api-channel-blob-storage-inject-fault-severe-rate

Inject fault in calls to blob storage for the channels.json blob using Simmy (severe injection rate).



ps-playback-api-channel-blob-storage-inject-fault-terrible-rate

Inject fault in calls to blob storage for the channels.json blob using Simmy (terrible injection rate).

CLIENT APP



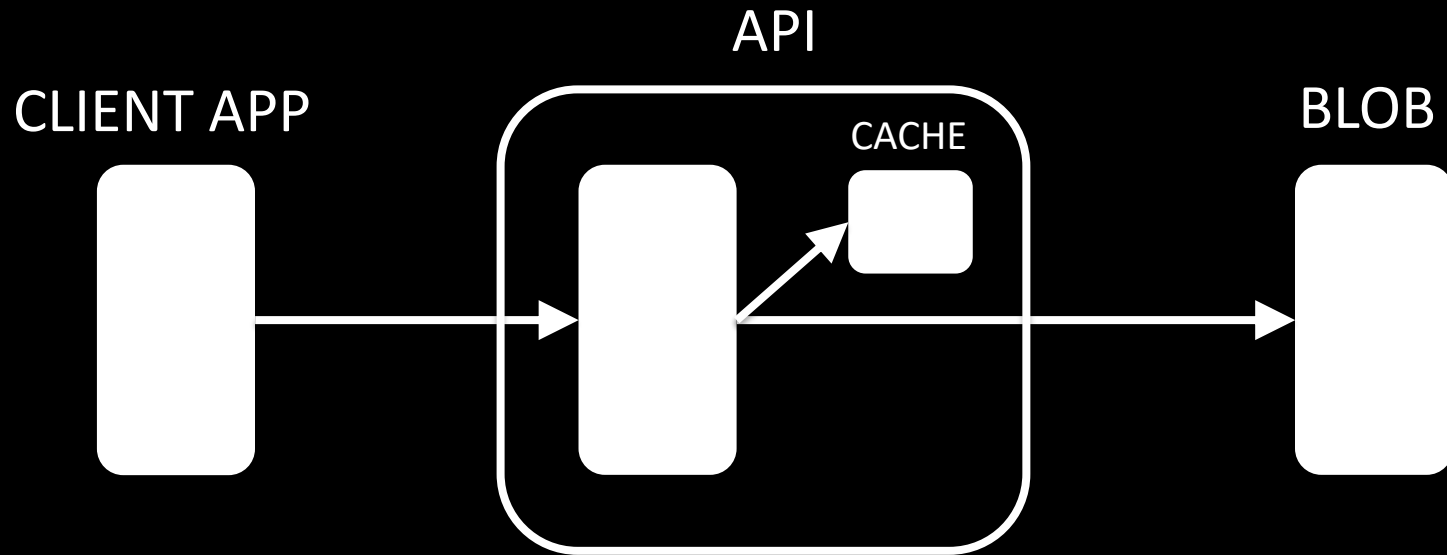
API
















BLOB

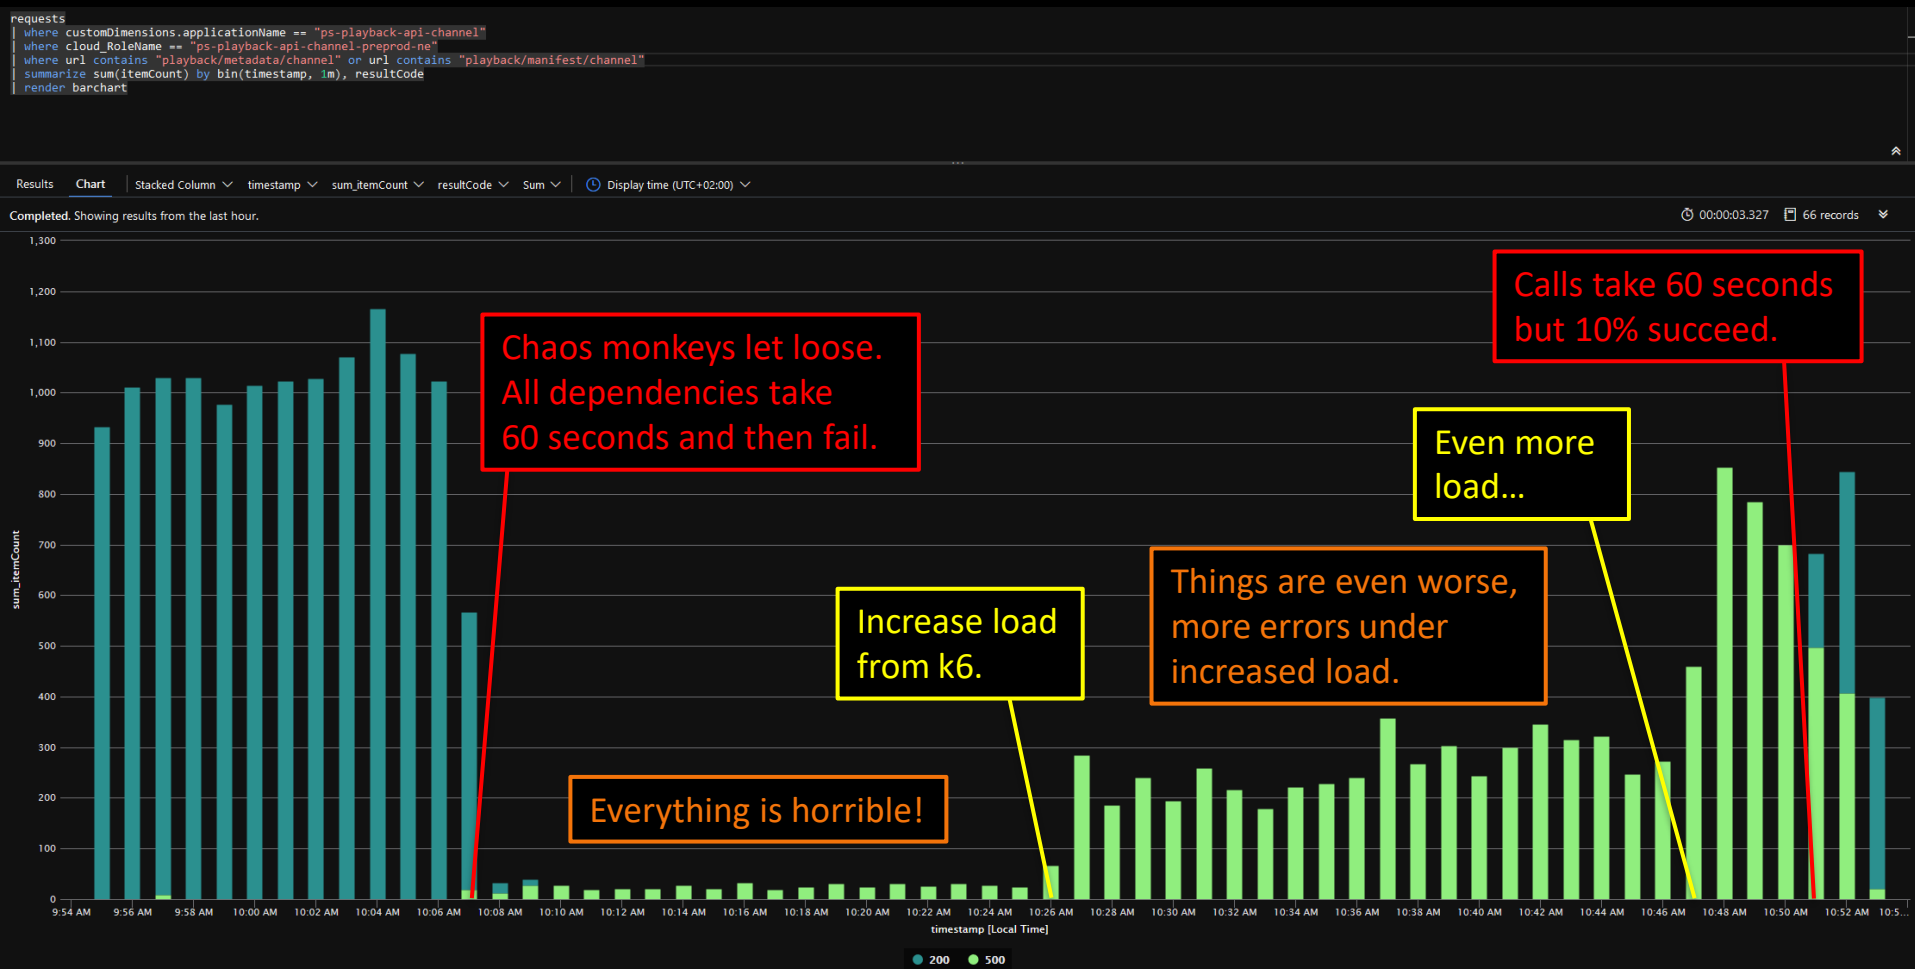


ORIGINAL DESIGN



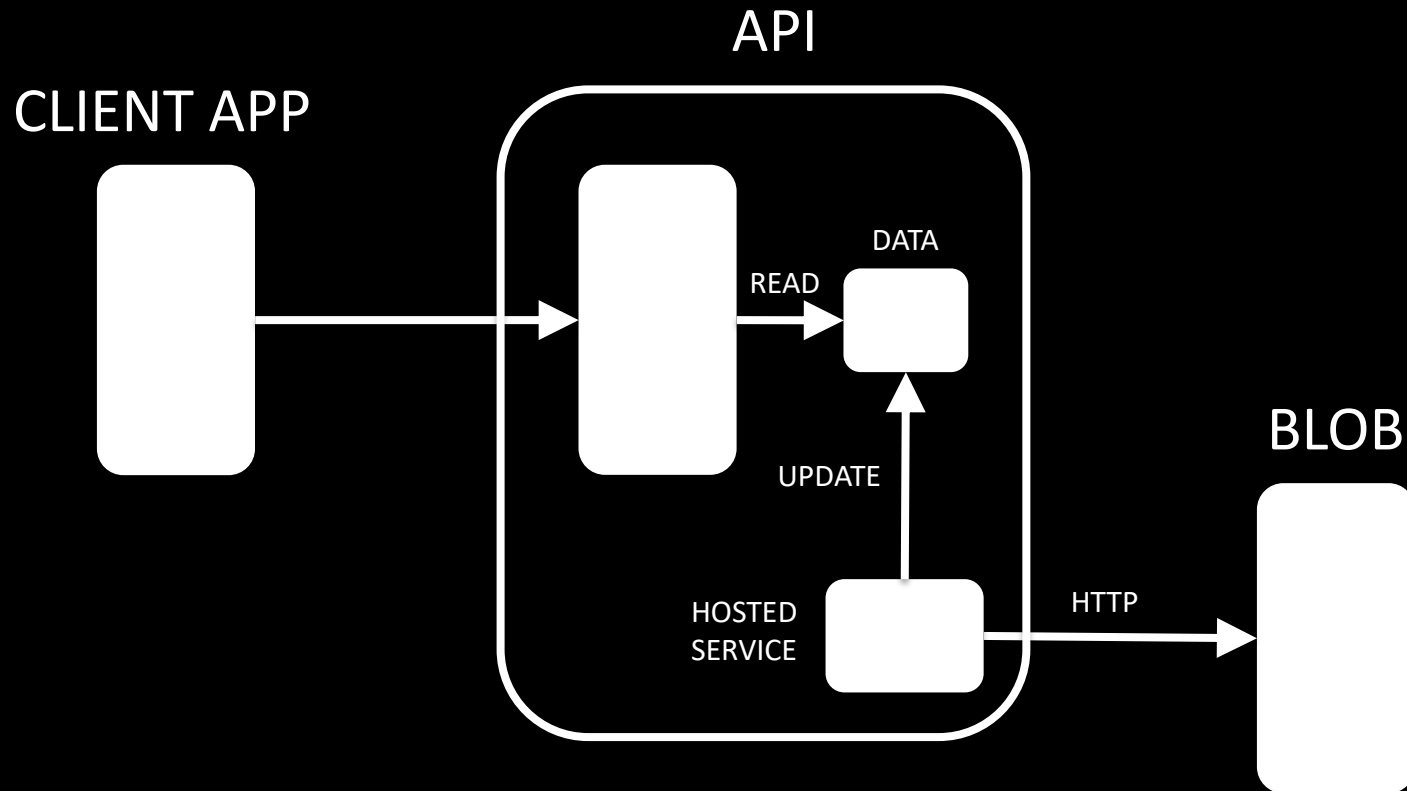
EXPERIMENT

	<input checked="" type="checkbox"/>	ps-playback-api-channel-blob-storage-inject-fault-always Inject fault in all calls to blob storage for the channels.json blob using Simmy.	default
	<input type="checkbox"/>	ps-playback-api-channel-blob-storage-inject-fault-low-rate Inject fault in calls to blob storage for the channels.json blob using Simmy (low injection rate).	default
	<input type="checkbox"/>	ps-playback-api-channel-blob-storage-inject-fault-moderate-rate Inject fault in calls to blob storage for the channels.json blob using Simmy (moderate injection rate).	default
	<input type="checkbox"/>	ps-playback-api-channel-blob-storage-inject-fault-severe-rate Inject fault in calls to blob storage for the channels.json blob using Simmy (severe injection rate).	default
	<input type="checkbox"/>	ps-playback-api-channel-blob-storage-inject-fault-terrible-rate Inject fault in calls to blob storage for the channels.json blob using Simmy (terrible injection rate).	default
	<input checked="" type="checkbox"/>	ps-playback-api-channel-blob-storage-inject-latency-always Inject latency in all calls to blob storage for the channels.json blob using Simmy.	default
	<input type="checkbox"/>	ps-playback-api-channel-blob-storage-inject-latency-low-rate Inject latency in calls to blob storage for the channels.json blob using Simmy (low injection rate).	default
	<input type="checkbox"/>	ps-playback-api-channel-blob-storage-inject-latency-moderate-rate Inject latency in calls to blob storage for the channels.json blob using Simmy (moderate injection rate).	default
	<input type="checkbox"/>	ps-playback-api-channel-blob-storage-inject-latency-severe-rate Inject latency in calls to blob storage for the channels.json blob using Simmy (severe injection rate).	default
	<input type="checkbox"/>	ps-playback-api-channel-blob-storage-inject-latency-terrible-rate Inject latency in calls to blob storage for the channels.json blob using Simmy (terrible injection rate).	default
	<input checked="" type="checkbox"/>	ps-playback-api-channel-blob-storage-latency-long-duration Any latency injected in calls to blob storage for the channels.json blob should be of long duration.	default
	<input type="checkbox"/>	ps-playback-api-channel-blob-storage-latency-medium-duration Any latency injected in calls to blob storage for the channels.json blob should be of medium duration.	default
	<input type="checkbox"/>	ps-playback-api-channel-blob-storage-latency-short-duration Any latency injected in calls to blob storage for the channels.json blob should be of short duration.	default

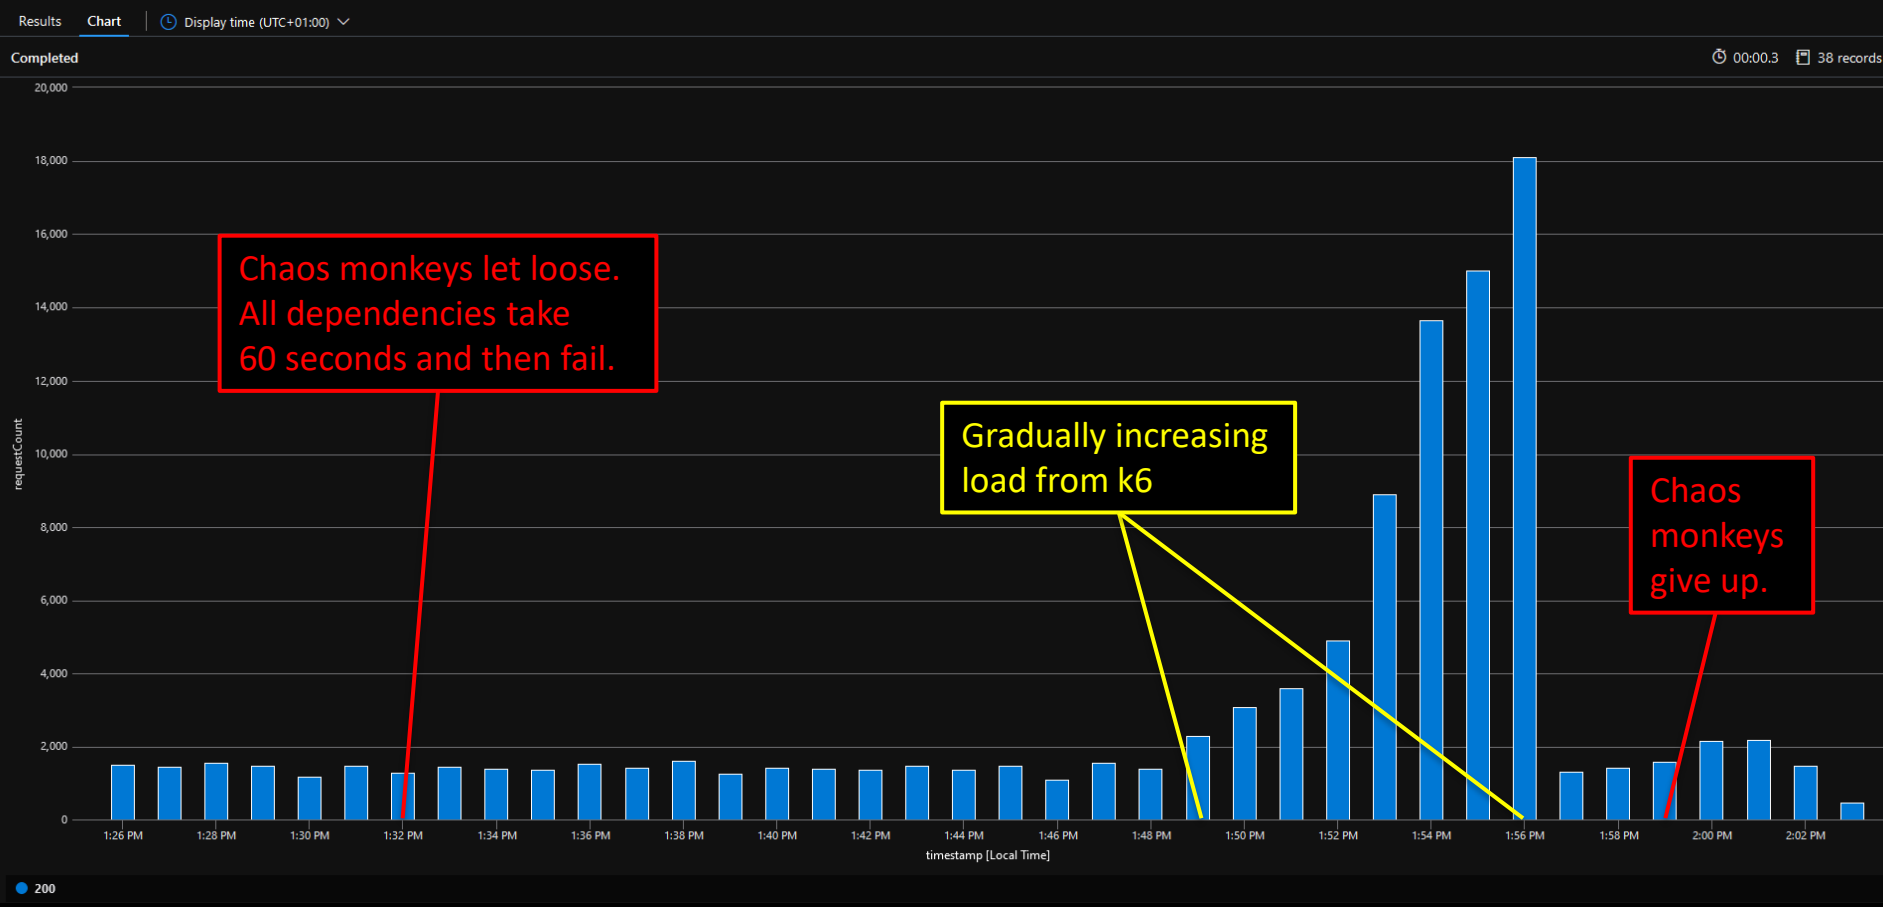


BUT STILL...

REVISED DESIGN



```
1 requests
2 | where timestamp > ago(46m)
3 | where customDimensions.applicationName == "ps-playback-api-channel"
4 | where cloud_RoleName == "ps-playback-api-channel-preprod-ne"
5 | where url contains "playback/metadata/channel" or url contains "playback/manifest/channel"
6 | summarize requestCount=sum(itemCount) by bin(timestamp, 1m), resultCode
7 | render columnchart
```



INJECTED FAULTS MADE IRRELEVANT

IS IT ROBUST?

TRICK QUESTION!

ROBUST ISN'T A BOOLEAN!

**VERY ROBUST AGAINST
BLOB DOWNLOAD PROBLEMS**

**NOT ROBUST AGAINST
AZURE GOING DOWN**

**THE CASE STUDY IS
A SIMPLE SCENARIO**

**MANY OTHER SCENARIOS
ARE MUCH HARDER**

**TRADE-OFFS CAN
BE MORE PAINFUL**

**ONLY AS GOOD AS
THE EXPERIMENTS**

**« I HAVEN'T BEEN ABLE
TO MAKE IT FAIL »**

SUMMARY

RELIABILITY ENGINEERING

REQUIRES FEEDBACK

**THE RIGHT MECHANISMS
ARE CONTEXT-DEPENDENT**

EXPLOIT CONTEXT TO FIND GOOD SOLUTIONS

ALWAYS VALIDATE

RUN EXPERIMENTS

LISTEN TO THE FEEDBACK

Monday

Room 6

09:00 - 17:00 (UTC+01)

2 Days

Building and testing resilient services over HTTP

Users' patience with services not working is gradually decreasing as the quality of services online is improving. With current trends of moving services to the cloud and building smaller and network-intensive services, meeting these expectations can be challenging for us developers. We want to be able to build services that we can run confidently despite partial failures and outages.

Venue: Rebel

.NET

In this workshop, you will learn how to simulate latency and failures in your web application and how to add strategies to deal with this. We will learn how changing the different parameters change the behavior of our application under load, and what trade-offs we ultimately must make. Because we build on HTTP and TCP/IP, we will also have fun sessions where we dig deeper into details in hands-on sessions to get a better understanding of the foundations on which we build.



Bjørn Einar Bjartnes

Bjørn Einar is a .NET developer and architect working at 4Subsea. He has a background from automation systems in the energy sector and has for the last 6 years worked for NRK TV as a backend developer and architect in the streaming service. His main topics of interest are domain driven design and functional programming. Keeping NRK TV's services up and running has been the main driver for his interest in resilient architecture. In the search for robustness, he has also learned to love HTTP.

[Linkedin](#)



Roger Hoem-Martinsen

Roger is a full-stack test developer with a varied background and a strong passion for building quality software solutions. In recent years, he has specialized in software testing. He started his career 15 years ago as a R&D engineer in signal processing, developing solutions for underwater communications. He has also worked with jet plane modems, gas detection systems, signal systems for railway, wireless condition monitoring systems and streaming services.

<https://github.com/bjartnes/bounded-disturbances>

